

# ECS404U: Computer Systems & Networks

Week 8: More into MIPS ISA:  
Memory Layout, Data Declaration, Branching  
instructions, I/O syscalls

---

Prof Edmund Robinson, Dr Arman Khouzani

November 09, 2020

EECS, QMUL

# Agenda/Objectives

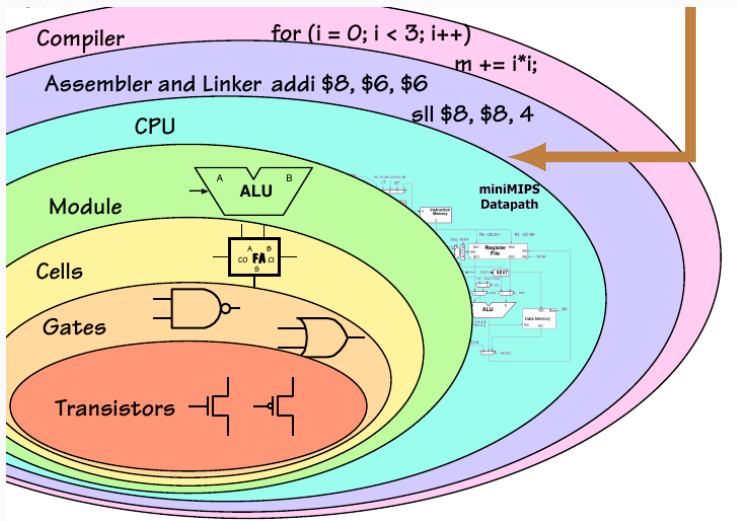
- ▶ Controlling the program flow: Sequencing (Jump and branching instructions)
- ▶ Understanding a simplified “memory layout” of a computer
  - ▷ Similarities and differences between instructions and data in the memory;
  - ▷ Declaring data in an assembly programme;
  - ▷ Memory Access (read and write);
  - ▷ Strings and Arrays;
- ▶ Simple I/O and syscalls
- ▶ Some practice MIPS assembly examples

## **A Quick Recap (where we are)**

---

# Machine Codes, Instruction Set Architecture

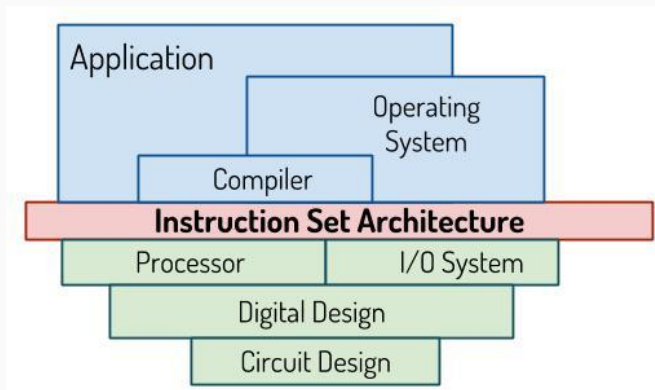
**ISA:** *the interface between hardware and software*



# Machine Codes, Instruction Set Architecture

- ▷ A given microprocessor has implemented a set of operations (e.g. addition, bitwise logics, etc.)
- ▷ These operations along with the available registers & modes of accessing memory (but without the detail of how they are implemented in hardware) constitute the **Instruction Set Architecture** of a processor.
- ▷ ISA provides an abstract computer for writing programmes that can directly run on the hardware.
- ▷ A **RISC** is an ISA that has a simpler and more efficient set of instructions (e.g. MIPS, ARM), unlike a **CISC** ISA (e.g. x86) which also involves complex instructions that take many clock cycles to perform.

# Machine Codes, Instruction Set Architecture



Instruction Set Architecture provides a simplified but sufficient description of the hardware interface, so that the programmes that are translated to it can run directly on the hardware.

# MIPS instruction set architecture

- ▷ The 32-bit MIPS machine instructions come in 3 formats: **R**, **I**, and **J**;
  - **R**: has references to **three registers**: e.g. some operation to be done on the content of two registers and the result be stored in the third register.
  - **I**: has reference to **two registers** and has an “**immediate**” value: e.g. some operation to be done on the content of one register and the immediate value, then the result be stored in the other register;
  - **J**: has no reference to any register, it only has an opcode and a **constant value**;
- ▷ In all three formats, the first 6 bits encodes the **opcode**, which specifies which operation to be done.

# MIPS instruction set architecture: Registers

- ▷ MIPS has **32 general-purpose registers**, each of which are **32 bits wide**;
- ▷ The first register (register **0** or **\$zero**) always has the **constant all zeros** (useful for moving data between registers, comparison instructions, etc);
- ▷ These 32 registers are indexed from **0** to **31**: in the 32-bit machine instruction, we designate which register through its number (with **5 bits**), but in the MIPS Assembly language, we also have **mnemonics** for them, e.g. register **0** is represented by **\$zero**, register **8** by **\$t0**, and so on.



# MIPS instruction set architecture: Registers

- ▶ Besides these general purpose registers, there are a few registers that are crucial for the working of the processor itself. Most notably:
  - ▷ **PC:** *Program Counter*, a (32-bit) register: contains the (memory) address of the next instruction to be fetched from memory;
  - ▷ **IR:** *Instruction Register*, a (32-bit) register: contains the most recently fetched instruction.
- ▷ The programmer does not have direct access to these registers. But then, how can we manipulate the flow of execution?

# Branching and Jumps

---

# Sequencing Instructions

The microprocessor fetches (from the main memory, i.e., RAM) the instruction whose address is determined by the content of the **PC** register.

- Normally, after execution of each instruction, the content of the **PC** gets incremented by **4**, to indicate the address of the next instruction in the memory;
  - ▶ *Question: why does PC get incremented by 4?*

# Sequencing Instructions

The microprocessor fetches (from the main memory, i.e., RAM) the instruction whose address is determined by the content of the **PC** register.

- Normally, after execution of each instruction, the content of the **PC** gets incremented by **4**, to indicate the address of the next instruction in the memory;
  - ▶ *Question: why does PC get incremented by 4?*
  - ▷ *Answer: in a 32-bit architecture, each instruction is 32 bits long, which is 4 bytes. Each byte of the memory has its own address, so the address of the (starting byte of the) next instruction is 4 plus the address of the (starting byte of the) current instruction.*


# Sequencing Instructions

- ▷ However, no interesting program follows this simple sequence (think e.g. “sorting”, or any “loop”)
  - We need to be able to control the flow of the program (e.g. based on the data, etc.);
- ▶ This is why we need “sequencing instructions” like jumps and branching.

# MIPS Sequencing Instructions: Jump

The simplest instruction to alter the default sequence of instructions is the **jump** instruction:

► **j** <label>



The J-format.

- ▷ *Effect:* loads the **PC** register with the address of the instruction that we have designated by a **label** (in our assembly code, using a **colon**).

# MIPS Sequencing Instructions: Jump

[details Not part of assessment!]

- ▶ ***Question:** Recall that each address is 32-bits (**PC** is a 32-bit register), but the **CONSTANT** in the **j** instruction is 26 bits long (it cannot be longer, why?). How do we get to a 32-bit address from a 26-bit constant?*
- ▷ ***Answer:** Two zeros are added to the right of it (why?) and for the remaining 4 bits, it uses the high order four bits (four leftmost bits) of whatever is in the PC!*
- ▶ ***Question:** Explain why using **j** instruction, you cannot jump to any 32-bit address in the memory. What range can you go to?*

# MIPS Sequencing Instructions: Branching

With only a jump instruction, you can at most have an infinite loop! But what about controlling the flow of the program based on some condition, i.e., **branching**? (Q: *incidentally, why is it called branching?*)



# MIPS Sequencing Instructions: Branching

With only a jump instruction, you can at most have an infinite loop! But what about controlling the flow of the program based on some condition, i.e., **branching**? (Q: *incidentally, why is it called branching?*)

- ▶ Branch-on-equal (branch-if-equal):

**beq**      **rs**, **rt**, **label**

- ▷ compares the content of registers **rs** and **rt**, and if they are equal, executes the sequence of instructions starting from the line designated by **label**.

# MIPS Sequencing Instructions: Branching


With only a jump instruction, you can at most have an infinite loop! But what about controlling the flow of the program based on some condition, i.e., **branching**? (Q: *incidentally, why is it called branching?*)

- ▶ Branch-on-equal (branch-if-equal):

**beq**     **rs**, **rt**, **label**

- ▷ compares the content of registers **rs** and **rt**, and if they are equal, executes the sequence of instructions starting from the line designated by **label**.
- ▶ There is also **bne**: branch-on-not-equal.


# MIPS Sequencing Instructions: Branching



The (16-bit) constant is treated as the *relative* address from our current instruction address. In particular, the effect is: if the content of register `rs` is equal to `rt`, then:

$$PC = PC + 4 * \text{CONSTANT}$$

# MIPS Sequencing Instructions: Branching




The (16-bit) constant is treated as the *relative* address from our current instruction address. In particular, the effect is: if the content of register **rs** is equal to **rt**, then:

$$PC = PC + 4 * \text{CONSTANT}$$

▷ Q: why multiply by 4?

# MIPS Sequencing Instructions: Branching



The (16-bit) constant is treated as the *relative* address from our current instruction address. In particular, the effect is: if the content of register `rs` is equal to `rt`, then:

$$PC = PC + 4 * \text{CONSTANT}$$

- ▷ Q: why multiply by 4?
- ▷ Q: how many instruction ahead/before can this go to?

# Memory Layout

---

# Memory Layout

One of the features of the von-Neumann architecture (as opposed to *Harvard* architecture) is that both data and instructions reside in the main memory.

# Memory Layout

One of the features of the von-Neumann architecture (as opposed to *Harvard* architecture) is that both data and instructions reside in the main memory.

- ▷ *Q: what is an advantage and a disadvantage of this feature?*



# Memory Layout

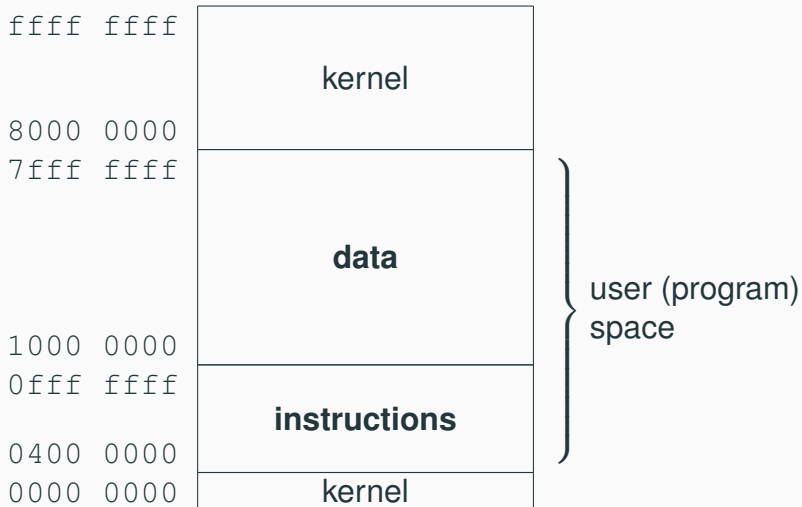
One of the features of the von-Neumann architecture (as opposed to *Harvard* architecture) is that both data and instructions reside in the main memory.

- ▷ *Q: what is an advantage and a disadvantage of this feature?*
- ▶ So we need to have a **memory “layout”** (a.k.a. a **memory “map”**) to keep the data and instructions in an orderly fashion, and separated as much as possible (why?!).

# Declaring Data

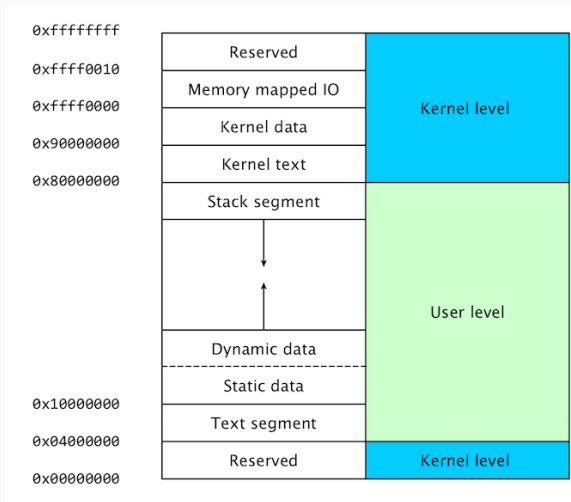
```
1 .text
2     li $v0, 10 # syscall to exit
3     syscall # Exit
4 .data
5     MSG:    .asciiz "abacus"
6     MY_INT: .word 15
7     MY_INT_ARRAY:
8         .word    1
9         .word    2
10        .word    3
11        .word    4
12        .word    5
13        .word    6
14        .word    7
15     MY_INT_ARRAY_LEN:    .word    7
```

# Memory Layout



Memory layout in a (32-bit) MIPS.

# Memory Layout



Detailed memory layout in a (32-bit) MIPS.

Recall that MIPS is an example of load/store architecture: any operation can only be done on the contents of registers. This means that in order to process data we have to first read them from memory into registers.

# Memory Access

We can read the data we have stored in memory using the following commands:

- Load address:

```
la    rt, <label of the data>
```

puts the address of the variable denoted by its label (which we have declared in the `.data` section of our code) in register `rt`.

- Load word:

```
lw    rt, offset(rs)
```

puts 4 bytes (a word) from the memory into register `rt`, whose starting address (address of the its first Byte) is `$rs + offset`.<sup>1</sup>

<sup>1</sup>Note: by putting a dollar sign \$ behind the name of a register, we mean the “content” of that register.

# Memory Access

To store values back in memory, we can use the following instruction:

- Store word:

**sw      *rt*,   *offset* (*rs*)**

the content of register ***rt*** is written in memory whose address (of its first Byte is)  **$\$rs + offset$** . Note that as the only exception to the general rule, here, the destination is written second!

# Memory Access

- ▶ Write a MIPS Assembly code that increment the value of **MY\_INT** that we have declared in the memory.



# Memory Access

- Write a MIPS Assembly code that increment the value of **MY\_INT** that we have declared in the memory.

```
1    la    $t0 , MY_INT      # t0=address(MY_INT)
2    lw    $t1 , 0($t0)      # t1=value(MY_INT)
3    addi  $t1 , $t1 , 1     # t1++
4    sw    $t1 , 0($t0)      # value(MY_INT)++
```

# **I/O: System calls**

---

# I/O: System Calls

Service	Code in <code>\$v0</code>	Arguments	Result
print integer	1	<code>\$a0</code> = integer to print	-
print string	4	<code>\$a0</code> = address of null-terminated string	-
read integer	5	-	<code>\$v0</code> holds integer read

Some of the Syscall services available in `QtSpim`, along with the detail of how to invoke them.

[Example demo.]

## Practice Example for MIPS Assembly

Suppose we have four variables initially loaded into registers `t0`, `t1`, `t2`, `t3`. Let's implement a *while loop* in MIPS Assembly:

```
while (t0 > t1) {  
    t2 = t2 + t0;  
    t0 = t0 - t3;  
}
```

Let's do this through steps. First, the easy part: the inside part of the loop:

<pre>add \$t2, \$t2, \$t0    # t2 ← t2+t0 sub \$t0, \$t0, \$t3    # t0 ← t0-t3</pre>
--

## Practice Example for MIPS Assembly

Let's turn it to a loop, and for now, without any conditions:

```
LOOP:
    # branching instruction to come here!
    add $t2, $t2, $t0    # t2 ← t2+t0
    sub $t0, $t0, $t3    # t0 ← t0-t3
    j  LOOP
EXIT:  # exit, or the rest of the code
```

So what should we put as the branching condition to exit from the loop?

# Practice Example for MIPS Assembly

The first (incorrect) attempt in the class was the following:

```
LOOP:
```

```
    beq $t0 , $t1 , EXIT    # goto exit if t0==t1
```

```
    add $t2 , $t2 , $t0     # t2 ← t2+t0
```

```
    sub $t0 , $t0 , $t3     # t0 ← t0-t3
```

```
    j LOOP
```

```
EXIT:    # exit , or the rest of the code
```

- although `$t0` keeps decrementing by `$t3`, it may never become equal to `$t1`: it may skip over it!
- even before entering the loop, we can have `$t0 < $t1`, then the **while** loop should have not executed at all!

Note that though this is wrong, this will get  $\sim 6/10$  mark!

# Practice Example for MIPS Assembly

The second (still incorrect) attempt was the following:

LOOP:

```
    slt $t4, $t0, $t1      # t4 ← (t0 < t1)
    bne $t4, $zero, EXIT   # if t4 != 0 goto EXIT
    add $t2, $t2, $t0      # t2 ← t2 + t0
    sub $t0, $t0, $t3      # t0 ← t0 - t3
    j LOOP
```

EXIT: # exit, or the rest of the code

- This implements **while** ( $t0 \geq t1$ ) instead!  
(because even if  $t0 == t1$ , the loop is executed)

Note that **slt** itself is not a branching instruction.

Again even though this is wrong, it'd get  $\sim 8/10$  mark.

# Practice Example for MIPS Assembly

And finally, a correct answer:

LOOP:

```
    slt $t4, $t1, $t0      # t4 ← (t1 < t0)
    beq $t4, $zero, EXIT   # if t4 == 0 goto EXIT
    add $t2, $t2, $t0      # t2 ← t2 + t0
    sub $t0, $t0, $t3      # t0 ← t0 - t3
    j LOOP
```

EXIT: # exit, or the rest of the code

Note that there may be more than one correct implementation, any correct answer will get 10/10 mark!

- More practice examples will be in your lab sheets 8, 9, 10, coursework, textbook, and relevant questions from previous year exams.